

# Final Project and Neural Networks COSC 425



THE UNIVERSITY OF  
TENNESSEE  
KNOXVILLE



# Announcements

- Lab 4 due on October 27
- No class on Thursday, October 26
  - Use that time to meet with your group/discuss your final project
- Final Projects:
  - Finalization of dataset to be used by **November 3**.
    - 5-point penalty if the dataset is not selected by that date.

**Questions about Lab 4?**

# Final Project


# Final Project

- Requirements for each of the final projects include:
  - Final project report: 6 to 8-page 2 column IEEE conference format report
  - In-class presentation: 3 minute in-class presentation overview of the project
  - Video presentation: 8–10-minute video presentation of the project
- Final projects will include:
  - The selection of a dataset
  - The analysis and visualization of this dataset
  - The application of machine learning approaches to the dataset
  - A discussion of hyperparameters required for this dataset

# Final Project Deadlines

- Team finalized by **October 19**.
- Finalization of dataset to be used by **November 3**.
  - 5-point penalty if the dataset is not selected by that date.
- Presentations will take place in class on **November 30 and December 5**.
- Final reports and video presentations will be submitted via Canvas and are due on **December 8**.

# Final Project





# Final Project

- You will be doing some data analysis on the dataset before you apply machine learning to it
  - Understanding and reporting the broader characteristics of the data
  - Cleaning the data
  - Visualizing the data (as best you can)
  - Preparing the data for machine learning approaches
    - Down-selecting
    - Scaling/transforming



# Datasets

- Sources:
  - US government data: <https://data.gov/>
  - World bank data: <https://data.worldbank.org/>
  - World Health Organization data: <https://www.who.int/data/gho/>
  - Miscellaneous data:
    - Kaggle: <https://www.kaggle.com/>
    - UCI ML Repository: <https://archive.ics.uci.edu/ml/index.php>
    - Azure open datasets: <https://docs.microsoft.com/en-us/azure/open-datasets/dataset-catalog>
- Finalization of dataset(s) to be used by **November 3**.
  - 5-point penalty if the dataset is not selected by that date.

# Machine Learning Approaches

- You will use machine learning approaches that we have covered in class.
  - **NOTE:** Your ML approach ***cannot*** be deep learning. We're covering that way too late and there is an entire class devoted to deep learning in this department.
- You can choose to use one or more machine learning approaches
- You will do an examination of the impact of hyperparameters on the machine learning approach(es)

# Project Report

# Project Report

- Your final project report should be structured as an academic paper for a conference or journal submission is structured.
  - Your final project report should use the IEEE 2-column conference format with the sigconf format.
  - There are both LaTeX and Word versions of this format
  - I recommend using Overleaf for LaTeX
- The final project report should be at least 6 pages and no more than 8 pages long.
- The report length includes any references, plots, and figures that you include.

# Project Report Components

- Abstract
- Introduction and Motivation
- Dataset
- ML Approaches and Methodology
- Results
- Discussion, Conclusion, and Future Work
- Contributions of Team Members

# Abstract

- 150-250 words to briefly introduce your dataset, the key question/problem you're addressing with ML applied to the dataset, and a brief summary of your results.
- Someone should be able to read just the abstract of your report and have an idea of what you did and what the results were.

# Introduction and Motivation

- You should address the following questions in the introduction and motivation section of your report
  - What is the dataset or datasets you chose for your project and why did you choose it/them?
  - What is the overarching goal of you're trying to achieve with ML on this dataset?
  - (Briefly) Which ML approaches are you using and why did you choose those?
- You should include a discussion here of what the goals are of the project and how you're determining whether you are successfully addressing those goals.
- Your introduction should include a summary of the rest of the paper as well.



# Dataset

- In this section, you should describe in detail the dataset that you're using
- Use visualizations to illustrate characteristics of the dataset
- Describe what features you're using and why
- If you omitted any features that are usually part of this dataset, describe why you omitted them
- Describe any transformations or pre-processing you performed on the data (and why)
- You must include at least two plots describing aspects of the data in this section

# ML Approaches and Methodology

- In this section, you should describe the ML approaches you chose to apply to this dataset and why you chose those ML approaches
- You should describe what hyperparameters you will be investigating of those approaches
- You describe which metrics you will be reporting and how those metrics are calculated
- You should describe in detail the experimental setup you have defined.
  - Are you comparing two different algorithms?
  - Are you comparing performance across hyperparameter values?
  - How are you defining whether the project is successful?

# Results

- In this section, you will describe the results of your approach.
- You should depict the results visually through plots, and you should provide a discussion of each plot and the results you obtained.
- Did something unexpected happen in the experiments?
  - It may be worthwhile to probe into that further to try to explain why it happened.
- Where appropriate, you should visualize the machine learning approach (through decision boundaries or some other visualization technique)
- If you can interpret the machine learning approach results, you should include a discussion of what you learned about the dataset from this interpretation (which features are most important, etc.)

# Discussion, Conclusion, Future Work

- You will provide a discussion of the results and any major conclusions you obtained
- You will also provide at least a paragraph of “future work” discussion.
  - If you had more time, what would you do next?
  - Did this project open up any new research questions?
  - Was there something else you would have liked to have done and didn’t get a chance to do?

# Contributions of Team Members

- This is a short section that is NOT included in the page count
- Here, you should describe what each team member did in the project
- There should be clear contributions of each member of the team
- Note: If you're working alone, you can omit this section

# **In-Class Presentation**

3-5 minutes – 4 slides

# In-Class Presentation

- You will have four TOTAL slides in your in-class presentation
  - Title Slide: Includes the title of the project and the names of the project team members (you should spend very little time showing this slide)
  - Dataset slide: Include a description and visualizations of the dataset
  - Machine learning methods slide: Include a description of the machine learning approach(es) you are applying and which hyperparameters you're examining
  - Results: Include an overview of the key results you obtained
- Note: You should keep text at a minimum and include visuals instead
- You will have a strict time limit to deliver the presentation
- All team members will stand at the front, but it is fine for one person to take the lead on presenting



# In-Class Presentation

- Most presentations will take place in class on December 5
- We will need at least 12 presentations to take place on **November 30**.
  - If you are willing to present on this day, please let me know via email
  - You will receive an automatic 5 points of extra credit on your final project for presenting early
- Slides for short presentations should be sent to Dr. Schuman by **November 29/December 4**

# Recorded Presentation

8-10 minutes

# Recorded Presentation

- You will create a video presentation of 8-10 minutes
- This should be a narration over a slide show
  - If you want to do it differently in some way, I'm open to that, but you need to get it approved first
- Every team member should present at least one part of the presentation
- You should submit this presentation as an mp4 or mov video file, and you will also submit the associated slides you presented as a PDF.
- I recommend recording this via Zoom.
  - Setup a Zoom session with the full group, have one person in charge of advancing slides and do a Zoom recording

# Rubric

- Report (70 points):
  - Abstract: 5 points
  - Introduction and Motivation: 10 points
  - Dataset: 10 points
  - ML Approaches and Methodology: 10 points
  - Results: 25 points
  - Discussion, Conclusions, and Future Work: 10 points
- In-Class Presentation (10 points)
  - Content (5 points)
  - Presentation quality (5 points)
- Recorded Presentation (20 points)
  - Slide content (10 points)
  - Presentation quality (10 points)

# Final Project Deadlines

- Team finalized by **October 19**.
- Finalization of dataset to be used by **November 3**.
  - 5-point penalty if the dataset is not selected by that date.
- Presentations will take place in class on **November 30 and December 5**.
- Final reports and video presentations will be submitted via Canvas and are due on **December 8**.

# What are you submitting?

- A zip file that includes:
  - PDF of your report
  - PDF of the slides presented in class
  - PDF of slides from pre-recorded presentation
  - Video file of pre-recorded presentation
  - Dataset (if the dataset is not publicly available)
  - Code files generated as a result of this project
- Each member of the team should submit the same zip file.

# Late Penalty

- There will be 10 points off per day late
  - Submitting at 12:00 AM on December 9 is considered a day late!
- Your final project can be submitted no later than December 12
- If your project is missing any of the required components, you will also receive a penalty if you submit those late



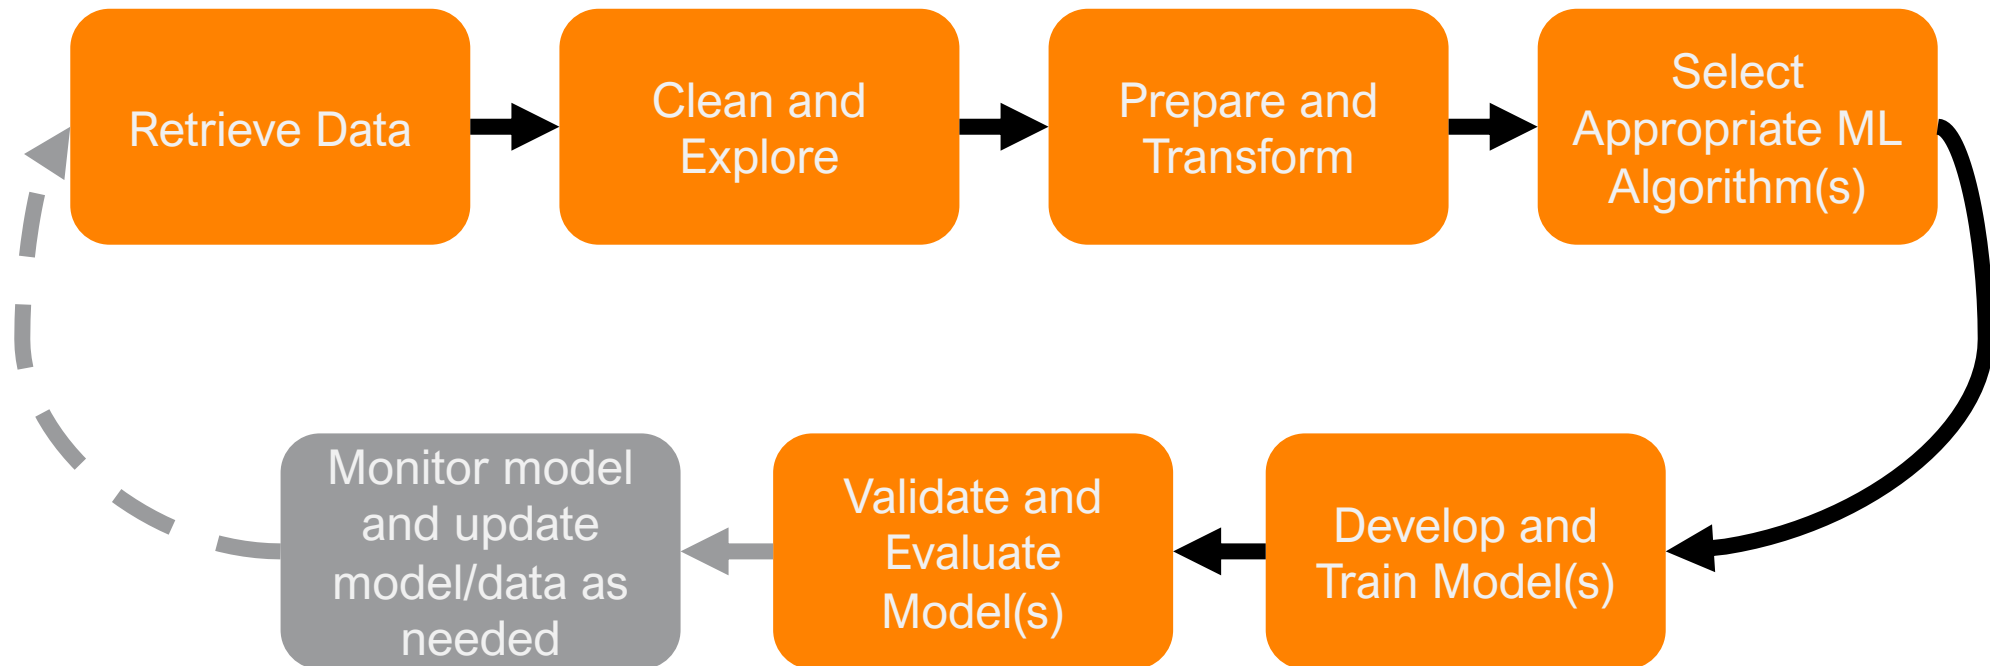
# Example Final Projects from 2022

- Compiler classification: GCC vs. Clang
- Spotify genre classification
- Stroke dataset classification
- Home price prediction
- Stock market prediction
- GTA Horse Betting

# Recommended Steps for Getting Started

1. Identify a domain you want to work in
2. Identify what types of problems you might solve in that domain (classification, regression, clustering)
3. Formulate the problem
  - Identify what the inputs/features will be
  - Identify what the outputs will be

# Final Project



# Pop Quiz

cs425


# Question

- My team and I have a good idea of what dataset/problem we're planning to work on
  - A) Yes
  - B) Sort of
  - C) No

# Multi-Layer Perceptrons

# Review on Neural Networks


- We learned about multi-layer perceptrons and showed they can deal with data that is not linearly separable
- We learned that a network a single hidden layer is a universal function approximator
- We showed how to train them with back-propagation






# Beyond Two Layers

- The definition of neural networks and back-propagation can be extended beyond a single hidden layer to any arbitrary directed acyclic graph (DAG)
- The structure is usually in layers, where each layer is fully connected to the next



# Beyond Two Layers

- However, all of the principles can be extended to any DAG
- In this example, we index our nodes in the graph as  $u, v$ .
- The activation before applying non-linearity at a node is  $a_u$  and after nonlinearity is  $h_u$
- The graph has D-many inputs who activations are given by an input example
- An edge  $(u, v)$  is from a parent to a child.
- Each edge has a weight  $w_{u,v}$
- $par(u)$  are all of the parents of  $u$



# Beyond Two Layers: Forward Propagation

---


**Algorithm 27** FORWARDPROPAGATION( $x$ )

---


```
1: for all input nodes  $u$  do
2:    $h_u \leftarrow$  corresponding feature of  $x$ 
3: end for
4: for all nodes  $v$  in the network whose parent's are computed do
5:    $a_v \leftarrow \sum_{u \in \text{par}(v)} w_{(u,v)} h_u$ 
6:    $h_v \leftarrow \tanh(a_v)$ 
7: end for
8: return  $a_y$ 
```

---


# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation



# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation







# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation




# Beyond Two Layers: Forward Propagation



# Beyond Two Layers: Forward Propagation



# Beyond Two Layers: Forward Propagation



# Beyond Two Layers: Back-Propagation

---

**Algorithm 28** BACKPROPAGATION( $x, y$ )

---

```
1: run FORWARDPROPAGATION( $x$ ) to compute activations
2:  $e_y \leftarrow y - a_y$  // compute overall network error
3: for all nodes  $v$  in the network whose error  $e_v$  is computed do
4:   for all  $u \in \text{par}(v)$  do
5:      $g_{u,v} \leftarrow -e_v h_u$  // compute gradient of this edge
6:      $e_u \leftarrow e_u + e_v w_{u,v} (1 - \tanh^2(a_u))$  // compute the “error” of the parent node
7:   end for
8: end for
9: return all gradients  $g_e$ 
```

---



# Wide or Deep?

Breadth vs. Depth

# Breadth vs. Depth

- Now, we know how to train both single hidden layer networks and arbitrary DAG networks
- We also know that single hidden layer networks are universal function approximators
- If single hidden networks are so great, why do we care about deeper (more layers) networks?

# Network Complexity


- The reason is that we want to keep the networks as simple as possible!
- We can show that there are functions that will require a huge number of hidden neurons if you force the network to be shallow (a single hidden layer) and only need a few neurons if you let it be deeper.

# Parity Function

- The parity function is a generalization of the XOR problem:

$$\text{parity}(x) = \sum_d x_d \bmod 2 = \begin{cases} 1 & \text{if the number of 1s in } x \text{ is odd} \\ 0 & \text{if the number of 1s in } x \text{ is even} \end{cases}$$

- We can easily define a circuit with depth  $O(\log_2 D)$  with  $O(D)$ -many gates for computing the parity function
- Each gate is an XOR gate
  - Organized in a complete binary tree
- You can do XOR with the network we went over in the first neural networks lecture



# Parity Function

- What does this mean?
- If we can go deep with our network, we can create a circuit that computes parity with a number of hidden units that is linear in its dimensionality
- Can you do this with a shallow circuit?
- NO!

# Parity Function Complexity

- Theorem: Any circuit of depth  $K < \log_2 D$  that computes the parity function of  $D$  input bits must contain  $O(e^D)$  gates.
- This is a famous result because it shows constant-depth circuits are less powerful than deep circuits
- It is generally believed that the same result holds for neural networks
- At the very least...it gives a strong indication that depth might be an important consideration in neural networks

# Impact of Depth on Parameters

- Heuristic: You need one or two examples for every parameter
- With this heuristic, a deep model could require exponentially fewer examples to train than a shallow model
- If deep is so much better, why isn't everyone using deep networks?
- (Actually, deep networks are now the norm, i.e., deep learning)

**Is deep always better?**



# Issues with Deep Networks



- Hyperparameter selection!
  - Deep networks have a lot more hyperparameters
- A single hidden layer network we only have to choose how many hidden units should appear in that layer
- With deep networks, we have to choose:
  - Number of layers
  - Number of units PER layer
- This can be daunting
- Hence, the field of neural architecture search

# Issues with Deep Networks

- As back-propagation makes its way backward through the model, the sizes of the gradients shrink
- If you are at the beginning of a deep network, changing one weight is unlikely to have a huge impact on the output, because it has to go through so many other neurons/units before it gets to the output
- So, the derivatives at the beginning of deep networks are small!
  - Back-propagation never really moves weights at the beginning of deep networks far from their random initialization
- This is called the vanishing gradients problem!

# Addressing these issues

- Obviously, we've overcome these issues with recent, massive, deep networks
- You can change the way initialization is done on a layer-by-layer basis, maybe in an unsupervised way
- You can use more sophisticated optimization beyond gradient descent
- Batch normalization can be used to help
- Other activation functions (like ReLU) that suffer less than activations like sigmoid
- Skip connections



# Basis Functions

# Neural Networks

- We know neural networks can mimic linear functions
- We also know they can learn more complex functions
- Can they learn to mimic a KNN classifier?
- Can they do it efficiently (i.e., without too many hidden units)?

# Swapping Activation Functions



- In a sigmoid network, we use:

$$h_i = \tanh(w_i \cdot x)$$

- We can change the activation function to instead be the radial basis function (RBF):


$$h_i = \exp[-\gamma_i \|w_i - x\|^2]$$

- Intuitively, these are little Gaussian bumps around the locations specified by the weight vectors, where  $\gamma_i$  specifies the width of the bump
- If  $\gamma_i$  is large, then only those that are really close to  $w_i$  have non-zero activations



# Training RBF Networks

- To train a Radial Basis Function (RBF) network, we have to find the Gaussian widths  $\gamma_i$ , the centers of the Gaussian bumps  $w_i$ , and the connections between the Gaussian bumps and the output unit  $v$ .
- We calculate  $v$  the same way, but our derivatives change for the other variable.



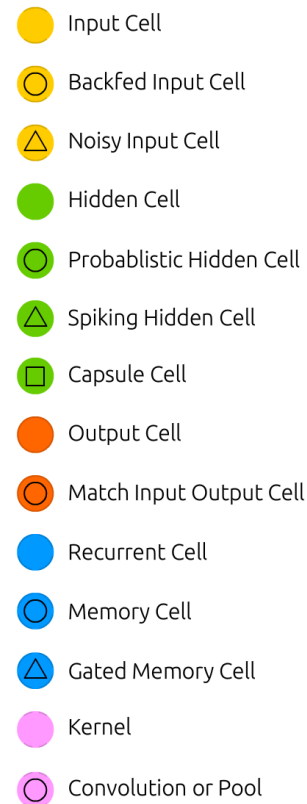
# RBF Networks

- Where should the Gaussian bumps be centered?
- You can apply back-propagation to find the weights OR you could specify ahead of time
- In particular, if you carefully choose  $\gamma$ s and  $\nu$ s, you can obtain something that actually looks a lot like distance-weighted KNN
- You can go further! Use back-propagation to learn good Gaussian widths ( $\gamma_i$ ) and voting factors ( $\nu$ ) for the nearest neighbors approach!



# We've barely scratched the surface


- Later in the semester we'll briefly cover deep learning
- There are lots of types of artificial neural networks and we definitely do not have time to cover them all in this class!




## A mostly complete chart of Neural Networks

©2019 Fjodor van Veen & Stefan Leijnen asimovinstitute.org


Perceptron (P)




Feed Forward (FF)




Radial Basis Network (RBF)




Deep Feed Forward (DFF)




Recurrent Neural Network (RNN)




Long / Short Term Memory (LSTM)




Gated Recurrent Unit (GRU)




Auto Encoder (AE)




Variational AE (VAE)



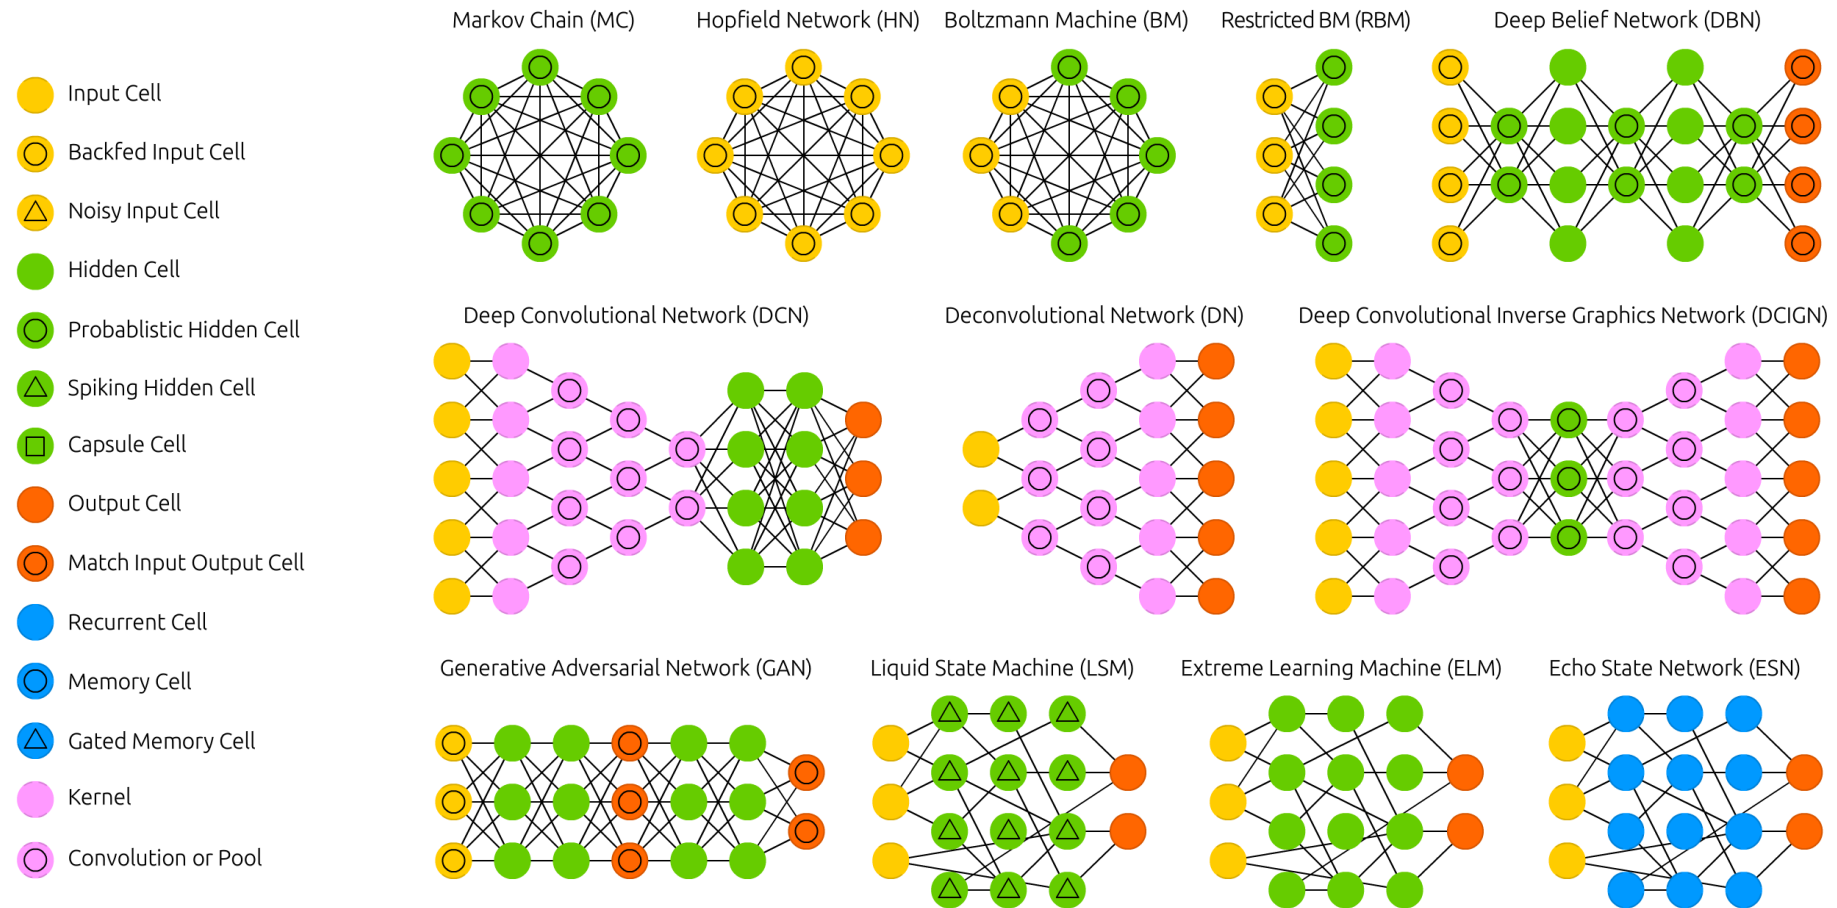
Denosing AE (DAE)

















Sparse AE (SAE)









# More Neural Networks



# More Neural Networks

-  Input Cell
-  Backfed Input Cell
-  Noisy Input Cell
-  Hidden Cell
-  Probabilistic Hidden Cell
-  Spiking Hidden Cell
-  Capsule Cell
-  Output Cell
-  Match Input Output Cell
-  Recurrent Cell
-  Memory Cell
-  Gated Memory Cell
-  Kernel
-  Convolution or Pool





# If you want to learn more about different types of neural networks...

- COSC 420/COSC 527: Biologically-Inspired Computing
  - Offered in spring
  - Lots of topics, but it includes Hopfield networks, liquid state machines, spiking neural networks
- COSC 525: Deep Learning
  - Offered in spring
  - Convolutional neural networks, autoencoders, generative adversarial networks, recurrent neural networks

# Announcements

- I will not have office hours tomorrow
- ***NO CLASS this Thursday, October 26***
- Lab 4 is due on Friday, October 27
- Final Projects:
  - Finalization of dataset to be used by **November 3**.
    - 5-point penalty if the dataset is not selected by that date.

# Announcements



Join us for EECS Mini Research/Internship/Career Fair!  
Speak to companies and professors in a more relaxed environment.  
Wearing jeans or business casual is recommended.  
Don't miss a great networking experience! (Dinner provided)

HOSTED BY SYSTERS: WOMEN IN EECS